

# Data Manipulation 2

STA 032: Gateway to data science Lecture 6

Jingwei Xiong

April 14, 2023

# Reminders

- HW 1 is due April 17th at 12pm
- HW 2 will be posted on the course website, due April 26 12pm.
- Please start the homework as soon as possible.
- Discussion will cover homework problems.

# Recap

- Data manipulation tools in tidyverse

- `select()`
- `arrange()`
- `slice()`
- `filter()`

Remember, before using all tidyverse functions, you need to `library(tidyverse)` first!

# Today

- Data manipulation tools continue
  - `distinct()`: filter for unique rows
  - `mutate()`: adds new variables
  - `count()`: create frequency tables
  - `summarise()`: perform column summarization operations
  - `group_by()`: for grouped operations
  - `pull()`: access column data as a vector or a number
  - `rename()`: rename an existing column
  - `inner_join()`, `left_join()`: join together a pair of data frames based on a variable present in both data frames that uniquely identify all observations
- Data visualization introduction

# Data: Hotel bookings

- Data from two hotels: one resort and one city hotel
- Observations: Each row represents a hotel booking
- Goal for original data collection: Development of prediction models to classify a hotel booking's likelihood to be cancelled ([Antonia et al., 2019](#))

```
hotels <- readr::read_csv("data/hotels.csv")
```

Source: [TidyTuesday](#)

# First question: What is in the data set?

.tiny[

```
dplyr::glimpse(hotels)
```

Rows: 119,390

Columns: 32

\$ hotel	Resort Hotel
\$ is_canceled	0, 0, 0, 0, 0
\$ lead_time	342, 737, 7,
\$ arrival_date_year	2015, 2015, 2
\$ arrival_date_month	"July", "July
\$ arrival_date_week_number	27, 27, 27, 2
\$ arrival_date_day_of_month	1, 1, 1, 1, 1
\$ stays_in_weekend_nights	0, 0, 0, 0, 0
\$ stays_in_week_nights	0, 0, 1, 1, 2
\$ adults	2, 2, 1, 1, 2
\$ children	0, 0, 0, 0, 0
\$ babies	0, 0, 0, 0, 0
\$ meal	"BB", "BB", "
\$ country	"PRT", "PRT",
\$ market_segment	"Direct", "Di
\$ distribution_channel	"Direct", "Di
\$ is_repeated_guest	0, 0, 0, 0, 0
\$ previous_cancellations	0, 0, 0, 0, 0
\$ previous_bookings_not_canceled	0, 0, 0, 0, 0

# distinct() to filter for unique rows

```
hotels %>%  
  distinct(market_segment)
```

```
# A tibble: 8 × 1  
  market_segment  
  <chr>  
1 Direct  
2 Corporate  
3 Online TA  
4 Offline TA/TO  
5 Complementary  
6 Groups  
7 Undefined  
8 Aviation
```

Recall: arrange() to order alphabetically

```
hotels %>%  
  distinct(market_segment) %>%  
  arrange(market_segment)
```

```
# A tibble: 8 × 1  
  market_segment  
  <chr>  
1 Aviation  
2 Complementary  
3 Corporate  
4 Direct  
5 Groups  
6 Offline TA/TO  
7 Online TA  
8 Undefined
```

## **distinct()** using more than one variable

```
hotels %>%  
  distinct(hotel, market_segment) %>%  
  arrange(hotel, market_segment)
```

```
# A tibble: 14 × 2  
  hotel      market_segment  
  <chr>      <chr>  
1 City Hotel Aviation  
2 City Hotel Complementary  
3 City Hotel Corporate  
4 City Hotel Direct  
5 City Hotel Groups  
6 City Hotel Offline TA/TO  
7 City Hotel Online TA  
8 City Hotel Undefined  
9 Resort Hotel Complementary  
10 Resort Hotel Corporate  
11 Resort Hotel Direct  
12 Resort Hotel Groups  
13 Resort Hotel Offline TA/TO  
14 Resort Hotel Online TA
```

**distinct()** is useful when you want to extract only the unique combinations of one or more columns in a data frame, and remove duplicate rows.



# mutate() to add a new variable

```
hotels %>%  
  mutate(little_ones = children + babies) %>%  
  select(children, babies, little_ones) %>%  
  arrange(desc(little_ones))
```

```
# A tibble: 119,390 × 3  
  children babies little_ones  
    <dbl>   <dbl>     <dbl>  
1         10      0         10  
2          0     10         10  
3          0      9          9  
4          2      1          3  
5          2      1          3  
6          2      1          3  
7          3      0          3  
8          2      1          3  
9          2      1          3  
10         3      0          3  
# ... with 119,380 more rows
```

What are these functions doing? How do to the same in base R?

Remember vector arithmetic? We can do similar things in homework 1 using mutate()

Remember vector arithmetic? We can do similar things in homework 1 using `mutate()`

---

HW1, Problem 4.1

---

HW1, Problem 5.1

HW1, Problem 7.5

```
temp <- c(35, 88, 42, 84, 81, 30)
city <- c("Beijing", "Lagos", "Paris", "Rio de Janeiro",
          "San Juan", "Toronto")
city_temps <- data.frame(name = city, temperature = temp)

city_temps %>% mutate(Celsius_temp = 5/9 * (temperature - 32))
```

	name	temperature	Celsius_temp
1	Beijing	35	1.666667
2	Lagos	88	31.111111
3	Paris	42	5.555556
4	Rio de Janeiro	84	28.888889
5	San Juan	81	27.222222
6	Toronto	30	-1.111111

Remember vector arithmetic? We can do similar things in homework 1 using `mutate()`

HW1, Problem 4.1

HW1, Problem 5.1

HW1, Problem 7.5

```
library(dslabs)
data(murders)
murders %>% mutate(rate = total/population * 100000) %>% head()
```

	state	abb	region	population	total	rate
1	Alabama	AL	South	4779736	135	2.824424
2	Alaska	AK	West	710231	19	2.675186
3	Arizona	AZ	West	6392017	232	3.629527
4	Arkansas	AR	South	2915918	93	3.189390
5	California	CA	West	37253956	1257	3.374138
6	Colorado	CO	West	5029196	65	1.292453

Remember vector arithmetic? We can do similar things in homework 1 using `mutate()`

HW1, Problem 4.1

HW1, Problem 5.1

HW1, Problem 7.5

```
murders %>% mutate(rank = rank(population)) %>%  
  arrange(desc(rank)) %>% head()
```

	state	abb	region	population	total	rank
1	California	CA	West	37253956	1257	51
2	Texas	TX	South	25145561	805	50
3	Florida	FL	South	19687653	669	49
4	New York	NY	Northeast	19378102	517	48
5	Illinois	IL	North Central	12830632	364	47
6	Pennsylvania	PA	Northeast	12702379	457	46

# count ( ) to create frequency tables

```
# alphabetical order by default
hotels %>%
  count(market_segment)
```

```
# A tibble: 8 × 2
  market_segment      n
  <chr>            <int>
1 Aviation          237
2 Complementary      743
3 Corporate          5295
4 Direct            12606
5 Groups            19811
6 Offline TA/T0     24219
7 Online TA          56477
8 Undefined           2
```

```
# descending frequency order
hotels %>%
  count(market_segment,
        sort = TRUE)
```

```
# A tibble: 8 × 2
  market_segment      n
  <chr>            <int>
1 Online TA          56477
2 Offline TA/T0     24219
3 Groups            19811
4 Direct            12606
5 Corporate          5295
6 Complementary      743
7 Aviation          237
8 Undefined           2
```

- Base R version: table()

# count() and arrange()

```
# ascending frequency order
hotels %>%
  count(market_segment) %>%
  arrange(n)
```

```
# A tibble: 8 × 2
  market_segment      n
  <chr>            <int>
1 Undefined          2
2 Aviation           237
3 Complementary      743
4 Corporate          5295
5 Direct            12606
6 Groups            19811
7 Offline TA/TO     24219
8 Online TA          56477
```

```
# descending frequency order
# just like adding sort = TRUE
hotels %>%
  count(market_segment) %>%
  arrange(desc(n))
```

```
# A tibble: 8 × 2
  market_segment      n
  <chr>            <int>
1 Online TA          56477
2 Offline TA/TO     24219
3 Groups            19811
4 Direct            12606
5 Corporate          5295
6 Complementary      743
7 Aviation           237
8 Undefined          2
```

# count ( ) for multiple variables

```
hotels %>%  
  count(hotel, market_segment)
```

```
# A tibble: 14 × 3  
  hotel      market_segment      n  
  <chr>      <chr>          <int>  
1 City Hotel Aviation          237  
2 City Hotel Complementary      542  
3 City Hotel Corporate        2986  
4 City Hotel Direct          6093  
5 City Hotel Groups        13975  
6 City Hotel Offline TA/T0    16747  
7 City Hotel Online TA       38748  
8 City Hotel Undefined         2  
9 Resort Hotel Complementary    201  
10 Resort Hotel Corporate      2309  
11 Resort Hotel Direct         6513  
12 Resort Hotel Groups        5836  
13 Resort Hotel Offline TA/T0   7472  
14 Resort Hotel Online TA     17729
```

# Order affects output when you count()

```
# hotel type first
hotels %>%
  count(hotel, market_segment)
```

```
# A tibble: 14 × 3
  hotel      market_segment      n
  <chr>      <chr>          <int>
1 City Hotel Aviation          237
2 City Hotel Complementary      542
3 City Hotel Corporate        2986
4 City Hotel Direct           6093
5 City Hotel Groups          13975
6 City Hotel Offline TA/TO    16747
7 City Hotel Online TA       38748
8 City Hotel Undefined         2
9 Resort Hotel Complementary    201
10 Resort Hotel Corporate      2309
11 Resort Hotel Direct         6513
12 Resort Hotel Groups         5836
13 Resort Hotel Offline TA/TO  7472
14 Resort Hotel Online TA     17729
```

```
# market segment first
hotels %>%
  count(market_segment, hotel)
```

```
# A tibble: 14 × 3
  market_segment hotel      n
  <chr>      <chr>          <int>
1 Aviation   City Hotel      237
2 Complementary City Hotel      542
3 Complementary Resort Hotel    201
4 Corporate  City Hotel      2986
5 Corporate  Resort Hotel    2309
6 Direct     City Hotel      6093
7 Direct     Resort Hotel    6513
8 Groups     City Hotel     13975
9 Groups     Resort Hotel    5836
10 Offline TA/TO City Hotel     16747
11 Offline TA/TO Resort Hotel    7472
12 Online TA  City Hotel     38748
13 Online TA  Resort Hotel   17729
14 Undefined  City Hotel         2
```



# summarize() for summary stats

```
# mean average daily rate for all bookings  
hotels %>%  
  summarize(mean_adr = mean(adr))
```

```
# A tibble: 1 × 1  
  mean_adr  
    <dbl>  
1    102.
```

- summarize() changes the data frame entirely
- Rows are collapsed into a single summary statistic
- Columns that are irrelevant to the calculation are removed

# summarize() is often used with group\_by()

- For grouped operations
- There are two types of hotel, city and resort hotels
- We want the mean daily rate for bookings at city vs. resort hotels

```
hotels %>%  
  group_by(hotel) %>%  
  summarize(mean_adr = mean(adr))
```

```
# A tibble: 2 × 2  
  hotel      mean_adr  
  <chr>      <dbl>  
1 City Hotel    105.  
2 Resort Hotel  95.0
```

- group\_by() can be used with more than one group

# Multiple summary statistics

`summarize` can be used for multiple summary statistics as well.

```
hotels %>%  
  summarize(  
    n = n(), # frequencies  
    min_adr = min(adr),  
    mean_adr = mean(adr),  
    median_adr = median(adr),  
    max_adr = max(adr)  
  )
```

```
# A tibble: 1 × 5  
      n min_adr mean_adr median_adr max_adr  
  <int>   <dbl>   <dbl>     <dbl>   <dbl>  
1 119390   -6.38    102.      94.6    5400
```

## `pull()`: access column data as a vector or a number

After the summarize result, it is a data frame (or tibble). not a vector or number.

```
# mean average daily rate for all bookings
hotels %>%
  summarize(mean_adr = mean(adr))
```

```
# A tibble: 1 × 1
  mean_adr
    <dbl>
1    102.
```

If we want to access the number, we can use the `pull()`

```
hotels %>%
  summarize(mean_adr = mean(adr)) %>% pull(mean_adr)
```

```
[1] 101.8311
```

## Another example:

```
hotels %>%  
  group_by(hotel) %>%  
  summarize(mean_adr = mean(adr)) %>% pull(mean_adr)
```

```
[1] 105.30447  94.95293
```

This can be useful when you want to assign a variable based on the result you calculated from the tidyverse workflow.

```
mean_adr = hotels %>%  
  summarize(mean_adr = mean(adr)) %>% pull(mean_adr)
```

## rename(): rename an existing column

The syntax is `rename(new_name = old_name)`.

Here we rename hotel column into hotel\_name.

```
hotels %>% select(hotel:lead_time) %>%  
  rename(hotel_name = hotel) %>% head()
```

```
# A tibble: 6 × 3  
  hotel_name    is_canceled lead_time  
  <chr>          <dbl>     <dbl>  
1 Resort Hotel      0       342  
2 Resort Hotel      0       737  
3 Resort Hotel      0         7  
4 Resort Hotel      0        13  
5 Resort Hotel      0        14  
6 Resort Hotel      0        14
```

# join family

Dplyr has a powerful group of join operations, which join together a pair of data frames based on a variable or set of variables present in both data frames that uniquely identify all observations. These variables are called keys.

- `inner_join`: Only the rows with keys present in both datasets will be joined together.
- `left_join`: Keeps all the rows from the first dataset, regardless of whether in second dataset, and joins the rows of the second that have keys in the first.
- `right_join`: Keeps all the rows from the second dataset, regardless of whether in first dataset, and joins the rows of the first that have keys in the second.
- `full_join`: Keeps all rows in both datasets. Rows without matching keys will have NA values for those variables from the other dataset.

# join family

To practice with the join functions, we can use a couple of built-in R datasets.

Dataset	Inner join	Left join	Right join	Full join
---------	------------	-----------	------------	-----------

```
data(band_instruments2)
head(band_instruments2)
```

```
# A tibble: 3 × 2
  artist plays
  <chr>   <chr>
1 John   guitar
2 Paul   bass
3 Keith  guitar
```

```
data(band_members)
head(band_members)
```

```
# A tibble: 3 × 2
  name  band
  <chr> <chr>
1 Mick  Stones
2 John  Beatles
3 Paul  Beatles
```



# join family

To practice with the join functions, we can use a couple of built-in R datasets.

Dataset	Inner join	Left join	Right join	Full join
---------	------------	-----------	------------	-----------

```
# Inner join  
band_members %>% inner_join(band_instruments2, by = join_by(name == a
```

```
# A tibble: 2 × 3  
  name  band    plays  
  <chr> <chr>   <chr>  
1 John  Beatles guitar  
2 Paul  Beatles  bass
```

# join family

To practice with the join functions, we can use a couple of built-in R datasets.

Dataset	Inner join	Left join	Right join	Full join
---------	------------	-----------	------------	-----------

```
# Left join  
band_members %>% left_join(band_instruments2, by = join_by(name == ar
```

```
# A tibble: 3 × 3  
  name  band    plays  
  <chr> <chr>   <chr>  
1 Mick  Stones <NA>  
2 John  Beatles guitar  
3 Paul  Beatles bass
```

# join family

To practice with the join functions, we can use a couple of built-in R datasets.

Dataset	Inner join	Left join	Right join	Full join
---------	------------	-----------	------------	-----------

```
# Right join  
band_members %>% right_join(band_instruments2, by = join_by(name == a
```

```
# A tibble: 3 × 3  
  name  band    plays  
  <chr> <chr>   <chr>  
1 John  Beatles guitar  
2 Paul  Beatles bass  
3 Keith <NA>    guitar
```

# join family

To practice with the join functions, we can use a couple of built-in R datasets.

Dataset	Inner join	Left join	Right join	Full join
---------	------------	-----------	------------	-----------

```
# Full join
band_members %>% full_join(band_instruments2, by = join_by(name == ar
```

```
# A tibble: 4 × 3
  name  band    plays
<chr> <chr>   <chr>
1 Mick  Stones  <NA>
2 John  Beatles guitar
3 Paul  Beatles bass
4 Keith <NA>    guitar
```

# Introduction to data visualization

- Why we need data visualization?

```
library(dslabs)  
data(murders)  
head(murders)
```

	state	abb	region	population	total
1	Alabama	AL	South	4779736	135
2	Alaska	AK	West	710231	19
3	Arizona	AZ	West	6392017	232
4	Arkansas	AR	South	2915918	93
5	California	CA	West	37253956	1257
6	Colorado	CO	West	5029196	65


- How is variable distributed?
- How can we identity patterns or relationships between variables

In contrast, the answer to all the questions above are readily available from examining this plot:

Picture

Code

explanation



In contrast, the answer to all the questions above are readily available from examining this plot:

Picture	Code	explanation
	<pre>library(tidyverse) library(ggthemes) library(ggrepel) library(ggplot2) r &lt;- murders  &gt;   summarize(pop=sum(population), tot=sum(total))  &gt;   mutate(rate = tot/pop*10^6)  &gt; pull(rate) murders  &gt; ggplot(aes(x = population/10^6, y = total, label = abb)) -   geom_abline(intercept = log10(r), lty=2, col="darkgrey") +   geom_point(aes(color=region), size = 3) +   geom_text_repel() +   scale_x_log10() +   scale_y_log10() +   xlab("Populations in millions (log scale)") +   ylab("Total number of murders (log scale)") +   ggtitle("US Gun Murders in 2010") +   scale_color_discrete(name="Region") +   theme_economist()</pre>	

In contrast, the answer to all the questions above are readily available from examining this plot:

Picture	Code	explanation
---------	------	-------------

Each state in the dataset was identified in this plot as a colored point with a label next to it. The total number of murders is shown on the y axis in log scale, and the populations are shown on the x axis in millions. The state name is indicated by the text label next to the points, and the color designates the state region. The average murder rate in the US was added as a gray line (in millions).



# Why we need data visualization

We are reminded of the saying "**a picture is worth a thousand words**". Data visualization provides a powerful way to communicate a data-driven finding.

Data visualization is the strongest tool of what we call *exploratory data analysis* (EDA). **John W. Tukey**, considered the father of EDA, once said,

"The greatest value of a picture is when it forces us to notice what we never expected to see."

Many widely used data analysis tools were initiated by discoveries made via EDA. EDA is perhaps the most important part of data analysis, yet it is one that is often overlooked.

The growing availability of informative datasets and software tools has led to increased reliance on **data visualizations** across many industries, academia, and government.

# Benefits of data visualization:

**Communication:** Data visualization provides a powerful way to communicate complex information to both technical and non-technical audiences.

**Exploration:** Data visualization allows us to explore data and identify patterns or trends that may not be apparent from numerical summaries alone.

**Identification of errors or outliers:** Data visualization can help us identify potential errors or outliers in our data that may impact our analysis.


**Hypothesis generation:** Data visualization can help generate new hypotheses or questions for further investigation.

# Another example

Picture

Code

explanation



# Another example

Picture	Code	explanation
---------	------	-------------

```
#knitr::include_graphics(file.path(img_path, "wsj-vaccines.png"))
data(us_contagious_diseases)
the_disease <- "Measles"
dat <- us_contagious_diseases |>
  filter(!state%in%c("Hawaii", "Alaska") & disease == the_disease) |>
  mutate(rate = count / population * 100000 * 52 / weeks_reporting) |>
  mutate(state = reorder(state, rate))
jet.colors <-
colorRampPalette(c("#F0FFFF", "cyan", "#00FFFF", "yellow", "#FFBF00", "orange", "red", "#7F0000"), bias
the_breaks <- seq(0, 4000, 1000)
dat |> ggplot(aes(year, state, fill = rate)) +
  geom_tile(color = "white", size=0.35) +
  scale_x_continuous(expand=c(0,0)) +
  scale_fill_gradientn(colors = jet.colors(16), na.value = 'white',
    breaks = the_breaks,
    labels = paste0(round(the_breaks/1000), "k"),
    limits = range(the_breaks),
    name = "") +
  geom_vline(xintercept=1963, col = "black") +
  theme_minimal() +
  theme(panel.grid = element_blank()) +
  coord_cartesian(clip = 'off') +
  ggtitle(the_disease) +
  ylab("") +
  xlab("") +
  theme(legend.position = "bottom", text = element_text(size = 8)) +
  annotate(geom = "text", x = 1963, y = 50.5, label = "Vaccine introduced", size = 3, hjust=0)
```

# Another example


Picture	Code	explanation
---------	------	-------------

A particularly effective example is a [Wall Street Journal article](#) showing data related to the impact of vaccines on battling infectious diseases.

One of the graphs shows measles cases by US state through the years with a vertical line demonstrating when the vaccine was introduced.

The plot shows the incidence rate of Measles in US states over time (years on the x-axis), represented by colored tiles for each state (on the y-axis). The incidence rate is calculated as the number of cases per 100,000 population per week, averaged over 52 weeks and adjusted for the number of weeks reporting data. States are sorted by their incidence rates, from lowest to highest, and are colored according to a gradient color scale, ranging from blue (low incidence rates) to red (high incidence rates). The plot includes a vertical line indicating the year when the Measles vaccine was introduced (1963). The plot is useful for visualizing how Measles incidence rates varied across US states over time, and how the introduction of the vaccine impacted the incidence rates.

In the talks [NewInsights on Poverty](#), Hans Rosling forces us to notice the unexpected with a series of plots related to world health and economics.




# Data visualization using ggplot2

Slide

Words 1

Words 2



- ggplot2 is the tidyverse's data visualization package
- create relatively **complex** and **aesthetically pleasing** plots
- syntax is **intuitive** and comparatively easy to remember.
- gg in "ggplot2" stands for Grammar of Graphics
- Inspired by the book Grammar of Graphics by Leland Wilkinson

# Data visualization using ggplot2

Slide

Words 1

Words 2

Throughout the lecture, we will be creating plots using the **ggplot2**<sup>[<https://ggplot2.tidyverse.org/>]</sup> package.

Many other approaches are available for creating plots in R. We chose to use **ggplot2** because it breaks plots into components in a way that permits beginners to create relatively **complex** and **aesthetically pleasing** plots using syntax that is **intuitive** and comparatively easy to remember.

One reason **ggplot2** is generally more intuitive for beginners is that it uses a grammar of graphics<sup>[<http://www.springer.com/us/book/9780387245447>]</sup>, the *gg* in **ggplot2**. This is analogous to the way learning grammar can help a beginner construct hundreds of different sentences by learning just a handful of verbs, nouns and adjectives without having to memorize each specific sentence. Similarly, by learning a handful of **ggplot2** building blocks and its grammar, you will be able to create hundreds of different plots.



# Data visualization using ggplot2

Slide

Words 1

Words 2

Another reason **ggplot2** is easy for beginners is that it is possible to create informative and elegant graphs with relatively simple and readable code.

To use **ggplot2** you will have to learn several functions and arguments. These commands may be hard to memorize, but you can always return back to this tutorial and grab the code you want. Or you can simply perform an internet search for **ggplot2 cheat sheet**.

# Next week:

Learn how to use `ggplot2` to generate the first example!

# Readings

- Chapter 4: The tidyverse
- Data Wrangling with Tidyverse
- Chapter 7: Introduction to data visualization